

Ballast: An LLM Database Steward for Vibe-Coded Applications

Richard Tibbetts
Bulrush Labs
Cambridge, MA, USA
tibbetts@gmail.com

ABSTRACT

Artificial Intelligence is pressuring software development to be faster and cheaper. The promise of agentic coding tools is that a non-developer can create applications, iteratively defining functionality with the agent making incremental changes. For professional developers, agentic tools accelerate development and tighten iteration cycles. A major obstacle to rapid development is managing schema evolution and migrating live data in an already-running application safely. The existing tooling for schema migration was designed for slower-moving systems with professional oversight. Today, projects built with AI ship schema change more than three times as often. Developers don't have time to validate all their migrations, and non-developers don't have the skills.

LLM-based coding agents like Claude Code can detect destructive migrations. Unfortunately, they can also convince themselves the migration will be OK, and execute it without human oversight. This has already produced high-profile incidents of data loss [20]. We show that *verdict-downgrading*, where the agent names a known-destructive operation and then permits it anyway, is a persistent flaw in coding agents. We present Ballast, a database steward that acts as the missing reviewer for schema changes and data migrations, blocking unsafe changes and alerting the user in clear natural language. Ballast inspects four aspects of a proposed change: potential for data loss, applicability to existing data, operational disruption, and fit to preexisting schema. The data loss checks, the most rigorously evaluated, enable Claude Sonnet 4.6 to give a correct verdict on every destructive fixture across five migration toolchains, block 93–98% of irreversible operations across 321 real agent-authored migrations from the AIDev corpus [10], with one confirmed false block in 150 clean agent-authored migrations (22 raw flags, hand-reviewed). Ballast is suitable for integration into data infrastructure itself.

1 INTRODUCTION

Computer programming is the biggest success story for large language models. Models are very good at producing working software. This has accelerated professional developers and also enabled non-developers to create systems by simply articulating specifications, known as “vibe-coding.” The promise is cheaper development, faster iteration cycles, and more customized software. In practice, agents frequently get into trouble operating on persistent data in production systems.

Studies of software development with agents show a steady decline in how much supervision people give to models: after agent adoption, pull requests receiving any human review fell from 89% to 68% and substantive review from 39% to 21%, with approval rates rising as inline comments fall [8, 28]. Unfortunately, there are cases

where a vibe-coded application goes off the rails, destroying persistent data or misleading human operators about what has happened to their data: in a widely reported 2025 incident, a coding agent deleted a production database during an explicit code freeze, then produced false status reports and false claims about recoverability [20]. A 2026 incident reportedly repeated the pattern on a routine staging task: the agent decided unprompted to delete a production database volume, taking the volume-level backups with it, despite a standing system-prompt rule against destructive irreversible commands. The rule was in the agent's context but was not enforced [2]. Risk of data loss slows adoption of these coding tools in higher-trust settings like financial and healthcare applications, domains that could benefit from rapid iteration if it could be combined with appropriate levels of trust.

In this paper we focus on a common style of application built by agents: database-backed web applications. Some are single-user, even throwaway one-shot code; but many are multi-user, even multi-tenant, and operate on persistent, valuable user data. Coding agents are very familiar with SQL databases, and many vibe-coding platforms provision Postgres by default [4, 11]. The last two decades produced rich tooling for managing schema evolution and data migrations. AI agents are capable users of this tooling, but they necessarily violate its key expectations: migration tools and linters were designed for a world where schema changes were infrequent, carefully considered, peer-reviewed, and executed through staging systems by professional database administrators or software operators.

Vibe-coded applications, by contrast, can produce dozens of small migrations in a week. Across 10,517 real vibe-coded repositories, the median project ran from first to last commit in under ten days, the fastest quartile at a median of 11.9 commits per day, over half of them on Supabase-provisioned Postgres [4]; migration commits specifically arrive 3.3× faster in agent-era projects (§2). Platforms apply these changes as part of the normal workflow, with no review step. Popular vibe-coding platform Lovable applies agent-written DDL to production Postgres with a chat approval as the only gate [11], and another platform, Replit, applies whatever schema changes the agent made in development to the production database when the app is published [17]. Schema changes happen incidentally, as a side effect of feature requests, and the person initiating the change is often unqualified, or short of time, to validate it. On Lovable's scaffold, 45% of repositories with migrations have at most two contributors, so the missing reviewer is a property of the project itself. To enable safe and rapid application development, we need a better understanding of database evolution and automated systems for keeping data safe.

In this paper we begin by detailing the observed shape of the problem, in published literature as well as newly collected and

```
-- add_multi_wallet_safe.sql  (Prisma, raw SQL)
DO $$ ... BEGIN
-- Check if walletAddress column exists
IF has_wallet_column THEN
CREATE TABLE IF NOT EXISTS "UserWallet" (...);
-- Migrate existing wallet addresses
INSERT INTO "UserWallet" (...)
SELECT ..., "walletAddress", "id", ...
FROM "User" WHERE "walletAddress"
IS NOT NULL ON CONFLICT DO NOTHING;
DROP TRIGGER IF EXISTS user_wallet_change
ON "User";
ALTER TABLE "User" DROP COLUMN
"walletAddress"; -- <- the anchor
END IF;
END $$;
```

Figure 1: Example of misleading comments and metadata (excerpted; the first line is our label, taken from the fixture’s name). The name says *safe*, possibly because of the guards, and the body appears to contain a data-preserving insert. But the `ON CONFLICT DO NOTHING` can lead to silent data loss, and the migration could be done in two steps, as Ballast recommends. The baseline Claude approved it in 9 of 10 trials. Ballast correctly blocked it in 10 of 10 trials.

synthetic data. We then present the Ballast system for conservatively validating the safety of database migrations, and analyze its accuracy against real-world changes, synthetic application histories, and adversarial evaluations. We survey complementary related work, and close with future work: these safety checks belong in data infrastructure itself, protecting data from the whims of an agentic coding harness.

2 DATABASE ERRORS BY AGENTS

Anthropic Claude, ChatGPT, and other state-of-the-art LLMs are quite capable of analyzing the safety of a database migration. But when such a model is asked to add a feature to an existing codebase, it looks for the shortest path to a good resolution. If the model encounters a destructive operation, it will tend to acknowledge the potential lack of safety, and then write a justification for making the change anyway. We call this safety failure **verdict-downgrading**. The built-in behavior of the model is to optimistically attempt to finish the task; that encourages the *appearance* of safety checks over actually blocking unsafe operations.

The same model may check that justification moments later, before the migration is committed and applied automatically to a production system. Agentic development thereby creates a circular source of trust: one model instance writes an optimistic filename, safety comment, or migration explanation, and a later instance treats that text as evidence without recognizing that it originated from the same fallible process now under review. Figure 1 shows the shape of the hazard.

In public repositories. A content-level analysis of 3,131 human-authored migrations across 24 open-source repositories finds **10.7%** unambiguously destructive; the same measurement on 3,305 agent-authored migration files from the AIDev agentic-PR corpus

[10] gives **9.7%**, similar to the human baseline [24]. What changes is exposure: agent-era projects ship schema change $3.3\times$ faster, with twice as many complex multi-table migrations (36.3% vs. 17.1%). Both are population effects, with no detectable authorship difference inside mixed repositories.

In agent traces. To measure the base rate of review error, we ran Claude Sonnet 4.6 as an unaided reviewer over recorded migrations and read its verdicts: one trial per migration on corpus sweeps, ten per prompt on the curated cases of §4. On the 205 AIDev migrations carrying an irreversible operation, it blocks **50.7%** and hedges a caution on **38.5%**. It names the hazard, then declines to block. On roughly three-quarters it voices concern in prose. Outright approvals are rarer (14 of the 205), and three of the fourteen acknowledge permanent data loss in the same rationale that approves; the purest writes “Data loss: Intentional — column and all values are deleted permanently;” “Rollback: None,” and then “Verdict: Safe to apply.” Another flatly denies the hazard: “No locking, data-loss, or NOT NULL concerns apply to a column removal,” written about a column removal. A condition has no force when no expert decides whether it was met; automation bias rubber-stamps the confident recommendation [15]. Even a machine-written warning does not move the verdict: Prisma’s generator writes “You are about to drop the column... All the data in the column will be lost” above the DROP, and the baseline blocks 2 of 10 trials with the warning present, 2 of 10 with it stripped; one rationale quotes the warning while concluding “Safe to apply.” Verdict-downgrading is a specific case of model behavior the literature already documents: sycophancy toward the material in front of the model [19], and stated conclusions that diverge from the model’s own reasoning [25].

In synthetic application histories. To measure the same behavior against known ground truth, we built ten multi-migration application projects (CRM, billing, IoT telemetry, e-commerce, ...) with AI agents: 131 migrations, each with a pre-registered verdict, across four toolchains. Reviewing them unaided, Sonnet 4.6 blocks 21 of the 25 destructive migrations and flags over a third of the clean ones, blocks included. Twenty-one of twenty-five is the baseline’s best showing on any corpus: synthetic migrations state what they do, while real destructive migrations arrive wrapped in additional context that legitimizes them. The baseline’s recall falls to 50.7% on real agent PRs and 17.5% on a production codebase (§4). One review asks the operator to confirm a column is empty before a DROP; with the fully automated deployment pipelines of §1, no one will read the request or verify the precondition.

3 THE BALLAST DATABASE STEWARD

Ballast is a database steward that uses LLMs to examine proposed changes. The prototype integrates with the agentic coding harness; future versions may deploy as part of database infrastructure, enabling a higher degree of security and safety (§6).

In agentic coding, an LLM like Claude Sonnet operates within a harness that controls its tools and system prompt. Many harnesses support configurable hooks and subagents. Subagents may share context with the main agent, but Ballast deliberately limits




Figure 2: Ballast prototype. The skill improves migrations during authoring; the validation agent reviews the final migration in fresh context, without the authoring session’s self-authored reassurance, and can block the apply.

that exposure: the validation agent decides based on the migration that will run, the schema, and prior source code, rather than the optimistic narrative of the main agent (Figure 2). Agent implementations fail with substantial common-mode correlation [18], so a same-family second look was never by itself the safety argument; Ballast’s authority comes from its rules. Early measurement of fresh-context review is encouraging but preliminary (§5).

The validation agent runs without shared context and tests whether a proposed migration is safe; if not, Ballast blocks it from running and notifies the human operator in domain language. Without the validation agent, any concerns the main coding agent raises stay hidden in thinking traces or voluminous output, subverted by possibly-flawed justification. The same rule set doubles as a *skill* in the authoring loop, yielding better, safer migrations whether or not the validation agent is in place.

Ballast rules are organized around four invariants for a safe migration: a recoverable downside, a clean fit with the data already in the tables, minimal operational disruption, and consistency with the schema that is already there. Verdicts are machine-parseable, carrying the decision with the object and rationale axis it applies to, so automation can grep for block while a human reads the prose explanation below it.

3.1 Preventing Data Loss

Ballast’s most rigorously evaluated rules are anchored on the destructive operation in the migration body (Figure 3): destructive operations are block-class *regardless of stated operational mitigations*. The principle is that a change is safe only if its downside is recoverable. Backups, staging runs, and low-traffic windows are real operational steps, but none of them makes a DROP reversible, so none of them may soften the verdict.

A block means “rewrite the safe way,” and the rule carries a rewrite recipe per operation. A column drop becomes the ignored-fields pattern: deploy code that no longer references the column, then drop in a later migration. A table drop becomes a rename to an archive name, with the real drop deferred past an archive window. Renames become expand-contract: add the new column, dual-write, move the readers, then drop the old one. A lossy type change becomes a new column and a validated backfill. Each recipe converts an irreversible step into stages that can be walked back.

For example Ballast blocks a real agent-authored migration that ships TRUNCATE automations_workflow under the comment

“Clean old test connections.” The comment may be true in development, where the table holds only leftover test rows; the migration file, though, runs in every deployment of the application. In a customer’s production database the same table holds live state. The steward’s objection:

“...a TRUNCATE in a versioned production migration file permanently destroys every row in that table with no way to recover short of a backup restore. Even if the table is expected to be empty in clean environments, it’s unsafe to ship a migration that silently wipes a table in production — any customer who has live automation workflows or pending connections will lose them without warning.”

3.2 Fitting the Existing Data

A structurally safe migration can still be wrong based on the rows it will affect. Ballast’s second set of rules judges a change against the state of the data: the audit queries, row counts, and cleanup lineage supplied with the request. Ballast blocks SET NOT NULL against a column whose audit reports nulls, a CHECK with existing violations, a foreign key added over orphaned rows, or a UNIQUE over duplicated data. The steward requires changes to be staged, cleaning the data before applying constraints.

Transform code gets the same scrutiny: does this backfill touch every row it claims to, and produce the values it intends? A backfill whose source expression can yield NULL collides with the SET NOT NULL that follows it. A narrowing cast overwrites values in place, unrecoverably. One subtle case we explicitly encode is the timezone-assumption conversion: casting a bare timestamp column with AT TIME ZONE 'America/New_York' when the writers

```
## Destructive ops are block-class
```

If the migration contains ANY of the following operations, emit `decision: block` regardless of stated operational mitigations (backup, staging, coordinate-with-app-deploy, run-during-low-traffic, verify-in-staging-first). Operational steps are real, but they do not transform a destructive DDL into a safe one for review purposes. A `block` verdict here means “the SQL needs to be rewritten as a safe-rewrite pattern before applying,” not “this can never ship.”

- DROP COLUMN / RemoveField
- DROP TABLE / DeleteModel
- Column and table renames
- ALTER COLUMN TYPE with visible evidence the new type is narrower [...]
- TRUNCATE

[...] If the migration contains a destructive op AND non-destructive ops in the same block, the block-class verdict still applies [...] Do not downgrade because most of the migration is fine.

Figure 3: The core rule of the Ballast skill, excerpted (the full skill is 1,749 words). The enumeration of mitigations that do not soften the verdict is the anti-downgrade mechanism.

stored UTC applies a silent offset to every historical row, and nothing fails.

We measured these rules in a pre-registered three-arm run of 480 trials with same-day arms. On the judged decision-quality scale of §4, they lift constraint introduction from 1.74 to 2.90 and data migration from 1.91 to 2.99, with the largest gains on the clean cases the baseline had failed. The same run measures the tension between these checks and the approve-side rules of §3.5.

3.3 Preventing Operational Disruption

A migration can be correct and recoverable and still take the application down while it runs. The availability rules flag operations whose lock windows scale with table size: `CREATE INDEX` without `CONCURRENTLY`, a constraint added without `NOT VALID`, a type change that rewrites storage, and the deceptive case, `ADD COLUMN...NOT NULL DEFAULT` with a volatile function, which forces a full table rewrite under an exclusive lock. Severity depends on a *declared* row count. Above ~10k rows, or on a table described as hot, the verdict is `block`; at small or unstated sizes the same pattern is a `caution` with the safe form attached. The declared count is the one hand-set constant in the system, and §4 shows it is where residual disagreement concentrates.

When operational context arrives with the request (table sizes, write rates, peak hours, scheduled jobs, a maintenance window), the steward gains a third verdict, `defer`, for a change that is safe to apply in the stated window and unsafe to apply during peak write load. A stated maintenance window is trusted as given and clears changes that would otherwise wait. The verdict exists because of how these applications deploy (§1): the platform applies migrations when the change lands, so timing judgment a DBA would exercise by feel has to be automated.

3.4 Conforming to Conventions

The fourth set of checks asks whether the change and the existing schema agree with each other. Given live introspection output alongside the committed migration history, the steward flags drift: objects that exist only in the database (an ad-hoc change applied by hand), migrations whose effects never arrived, and silent divergences where the migration text and the database disagree about a default or a type.

Convention guidance, we found, belongs in project context rather than in the steward. A generation-time skill telling the model how to scope tenant tables added nothing on a project whose context already stated the conventions. The steward checks conformance with what the project declares, and declines to dictate conventions of its own.

3.5 Reducing False Blocks: Approval Rules

A steward that blocks too eagerly gets ignored or uninstalled, and on clean input the model errs that way. On one canonical, correct policy prompt, 29 of 30 trials returned `caution` or `block` with the same invented defense-in-depth concern; those 30 span the unaided model and two early rule sets that predate the approve-side rules.

To reduce false blocks, Ballast includes approve-side rules. Mechanically they are paragraphs of prose in the same rule document

the model reads: bound concerns to the SQL and context in front of you, approve cleanly when nothing is wrong, and do not flag a change for what it could additionally enforce. Individual checks carry their own suppressors; the idempotency check, for example, stays quiet on framework-managed migrations, whose ledger already prevents re-application.

The two directions conflict. Given approve-side guidance alone, the model approves real hazards. An early approve-cleanly rule worsened decision quality on real destructive migrations from 0.64 to 0.36, of a possible 3. A variant carrying only the approve-side paragraphs considered a lossy narrowing cast, described the cent truncation it causes, and still softened the verdict to `caution` in 10 of 10 trials; the full rule set, whose checks name lossy casts, blocked all 10. Approve-side rules work only in tension with checks that name the hazards.

4 EVALUATING BALLAST

We evaluate Ballast against real-world changes (human- and agent-authored), synthetic application histories with pre-registered verdicts, and adversarial evaluations designed to probe edge cases. Table 1 collects the verdict counts by source; the paragraphs below give each source’s method and caveats. Every evaluation runs the reviewer over recorded migrations; no verdict in this paper was enforced against a live database. Where we score decisions rather than count verdicts, an LLM judge assigns a *decision quality* score (*dq*) from 0 to 3: 3 is the right verdict naming the specific concern, 0 the wrong verdict.

Real agent-authored migrations. We ran the reviewer over all 321 destructive migrations from the AIDev probe of §2, with ground truth derived from the DDL text itself. Recall tracks real irreversibility: counting firm blocks only, **97.8%** on `DROP COLUMN` and **93.3%** on `DROP TABLE` (an alter-column bucket is left unflagged by design). Every one of the 17 drop-class approves is defensible on reading the diff (state-only operations, drops in `down()`). Forward-history mining of the same 321 migrations found observable regret rare but real: one confirmed downstream failure (a type replacement whose queries began failing seventeen days later) and one project hand-executing the three-PR expand-contract Ballast recommends. False positives are counted on *clean* migrations only. The verdict-downgrading of §2 has a clean before and after (Table 1): where the unaided model hedges `caution` on 38.5% of the irreversible migrations, Ballast hedges on **zero** and blocks **94.1%**. Every Ballast caution sits in a recoverable or indeterminable class. The unaided model already voices concern in prose (§2); the skill converts that hedged prose into a firm, machine-parseable gate; in the frontmatter-scored comparison no case moved the other way.

A private production codebase. We ran Ballast against a stratified 129-migration sample of a real production Django codebase largely authored by coding agents. The sample’s 58 destructive-verb migrations include one state-only no-op, the skill’s own safe-rewrite idiom found in the wild; Ballast blocked all **57** genuinely destructive migrations and approved the no-op. The baseline firm-blocks 10 of the 57 and approves 39 outright. Fixtures let the baseline reach 84% recall; real history drops it to 17.5%,

Table 1: Ballast compared with the baseline model (no skill) across four sources of migrations. “Cases” counts the migrations reviewed from each source: the total, and the destructive subset scored (for AIDev, the 205 migrations carrying an irreversible operation). “Blocks” counts correct blocks of destructive migrations: higher is better. “False blocks” counts blocks of clean migrations: lower is better. Of Ballast’s 22 raw flags on 150 clean AIDev migrations, hand review confirmed only one as a true false block; the baseline’s AIDev figure is raw, and the remaining false-block figures are upper bounds that also count cautions. The curated-prompts row counts trials over 22 cases: nine cross-tool, seven adversarial fixtures, six real Prisma migrations ($N=10$ each). Ballast blocks nearly all destructive migrations from every source while rarely blocking clean ones.

source	cases		baseline		Ballast	
	total	destructive	blocks	false blocks	blocks	false blocks
agent PRs, public repos (AIDev)	471	205	50.7%	15.3%	94.1%	1/150 (22 raw)
private production codebase	129	57	10/57	7–11%	57/57	7–11%
synthetic application histories	131	25	21/25	36.7%	25/25	21.5%
curated prompts, 5 toolchains	22	20	–	–	200/200	0/20

because real destructive migrations arrive wrapped in legitimizing context. The probe also surfaced two verified bug catches in migrations that ran in production in 2023. A data-restore helper guard checks one field but reads another, silently stamping every touched row’s `updated_at` with its `created_at`; the same helper was copy-pasted into a second app and caught independently (3/3 trials, verified against source). The baseline catches this bug too; what Ballast adds here is the verdict discipline. The failure presented as in-the-wild studies predict: inside a plausible, well-commented, confidently written migration. Caveats: our own codebase (migrations predate this research programme), one model family, 129/573 sampled.

Synthetic application histories. On the synthetic corpus of §2, destructive recall is 25/25, stable under $4\times$ resampling (100/100 block verdicts). On the non-destructive migrations the baseline flags more and blocks harder (36.7% against Ballast’s caution-heavy 21.5%). Disagreement concentrates on the availability rule’s declared-row-count anchor (§3); its three block-expected misses all softened to caution, never approve, and *across 200 resample calls on this corpus, no forward destructive operation was ever approved*. In a full deployment of the first three projects, the skill twice overruled the answer key and was right, catching a self-recursive RLS policy and a non-deterministic backfill silently locked in by a following NOT NULL.

Other models. Cross-model, the core works smoothly: on GPT-5, destructive recall stays 12/12 in every arm, erring over-cautious on clean framework-wrapped migrations rather than missing destructive content.

Adversarial evaluations. Ballast was authored against Django content. Applied identically to Diesel, Rails, and Laravel destructive migrations, it blocked all 90 trials across the nine cross-tool prompts (Figure 4). Constructed counter-pulls did not trick it: virtue-named filenames over `remove_column` bodies, files literally named `_safe.sql` with `DROP`-class bodies, and dense in-body reassurance (guard blocks, `INSERT...SELECT` “preservation,” safety comments) that drags the *baseline* to 0.40 on the worst case while the rule stays at ceiling (Table 1). Figure 1 shows the sharpest case. Evidence spans five toolchains; in addition to the tools above we




Figure 4: Accuracy of the baseline model is dependent on exactly how migrations are written. Prisma’s reassurance-heavy raw SQL (Figure 1) confounds the baseline model, while Laravel’s explicit `dropColumn` is easier to understand. Ballast was originally written against Django but applies near perfectly across other systems, and blocks all destructive migrations. The Prisma pair shows the adversarial fixture of Figure 1; on the six real Prisma migrations of Table 1 the baseline scores 1.30 on destructive content.

also included Prisma migrations taken from a production repository. Across the six Prisma migrations, Ballast blocked all 40 destructive trials and approved all 20 clean ones, Prisma’s first measurement with clean cases included. Deliberate prompt injection into migration text remains untested.

5 RELATED WORK

Autonomous and LLM-assisted database management. Removing the human DBA has been a popular research topic for years: self-driving DBMSs [16], a decade of automatic tuning [26], LLM diagnosis [29]. These systems act to improve performance; Ballast earns its keep by *withholding* approval. The closest prior work is Horizon [6], which also pairs an LLM with hard programmatic checks in a migration setting. Horizon verifies that a *translated* SQL corpus is semantically equivalent across systems, whereas Ballast judges whether a *forward* schema-evolution migration is safe to apply to live state. The schema is changing; there is no reference to diff against.

Migration frameworks and linters. Two decades of tooling make migrations easier to author and keep environments consistent; linters [5, 9] verify static rules, online-DDL tools remove lock windows [14], and declarative tools diff a desired schema against the live database and compute the migration [3], which detects drift mechanically much as §3’s drift review does. None is designed to operate without human oversight, exercise judgment, or communicate with non-developers.

Constraints, snapshots, and branching. Database constraints protect data from a badly-behaved application; a badly-behaved migration can simply rewrite them. Branching and point-in-time recovery [13] are powerful recovery primitives, and no substitute for avoiding the bad migration in the first place.

Agent validation and the review gap. Insufficient human oversight of coding agents is extensively measured [23], including declining review under rising agent throughput [8, 28]. Cross-sectional agent-vs-human review comparisons are contested [1]; our measurement is *within* the agent population, destructive vs. non-destructive, which sidesteps that dispute. Agent-governance systems such as Progent [21] and IsolateGPT [27] gate tool calls through deterministic policy or isolation; Ballast adds database semantics those policies do not have.

Fresh-context review. Several practitioner systems review agent work in a fresh-context subagent without measuring the effect of the context isolation. AI Control [7] treats shared actor/checker identity as a first-class safety concern (for deliberate subversion rather than sycophancy); CriticGPT critiques read only the (question, answer) pair, with no author session [12]. One preprint measures the isolation directly: fresh-session review beat same-session self-review on injected errors, 28.6% vs. 24.6% F1 [22]; we do not measure it here.

6 FUTURE WORK

Data infrastructure integration. Ballast runs in the coding-agent environment and can offer only limited security guarantees: the main agent could disable Ballast or rewrite its rules to be less restrictive. Moving the final check into the data infrastructure would isolate it, enable secure audit trails, and give the steward direct access to production data. Live table access could improve detection.

Formal methods. Ballast judges migration code with an LLM; a sufficiently complex migration might sneak a bad operation past judgment. Theorem-proving-assisted validation of migration semantics is increasingly accessible and would harden the gate.

Other data-safety constraints. The architecture extends to read-path constraints, such as user privacy and multi-tenant isolation, and to “schemaless” stores.

7 CONCLUSION

Agentic coding is transforming software development: faster, cheaper, more accessible. To fully enjoy those benefits we need a new generation of database infrastructure that aligns with agent behavior: infrastructure that can protect itself from erroneous or dangerous operations undertaken by agents. Ballast demonstrates that a skill for guidance plus a fresh-context agent for validation

can block nearly all destructive operations, every destructive fixture across five toolchains and 93–98% of irreversible operations in the wild, with one confirmed false block in 150 clean agent-authored migrations (22 raw flags, hand-reviewed). When the database’s author is a language model, the steward does not have to be smarter than the author. It has to decide which of the author’s signals can be trusted, and block, advise, or ignore accordingly.

REFERENCES

- [1] Shyam Agarwal, Courtney Miller, Christian Kästner, and Bogdan Vasilescu. 2026. 3100 Opinions on Code Review in an AI World. *arXiv:2607.07980* [cs.SE]
- [2] AI Incident Database. 2026. Incident 1469: PocketOS Production Database Was Reportedly Deleted by Cursor AI Agent Running Claude Opus 4.6. Responsible AI Collaborative, <https://incidentdatabase.ai/cite/1469/>.
- [3] Ariga. 2026. Atlas: manage your database schema as code. Atlas Documentation, <https://atlasgo.io/docs>.
- [4] Junquan Deng, Zhiyu Fan, and Ruijie Meng. 2026. Understanding the (In)Security of Vibe-Coded Applications. *arXiv:2606.23130* (2026).
- [5] Steve Dignam and contributors. 2024. Squawk: a linter for PostgreSQL migrations. <https://github.com/sbdc/d/squawk>.
- [6] K. Venkatesh Emani, Wenjing Wang, Zi Ye, et al. 2025. Horizon: Robust Checks for SQL Migration Using LLMs. *Proceedings of the VLDB Endowment* 18, 12 (2025), 5259–5262. <https://doi.org/10.14778/3750601.3750646>
- [7] Ryan Greenblatt, Buck Shlegeris, Kshitij Sachan, and Fabien Roger. 2024. AI Control: Improving Safety Despite Intentional Subversion. *arXiv:2312.06942*.
- [8] Hao He, Shyam Agarwal, Yegor Denisov-Blanch, et al. 2026. AI Writes Faster Than Humans Can Review: A Longitudinal Study of an Enterprise 2x Mandate. *arXiv:2607.01904* [cs.SE]
- [9] Andrew Kane. 2024. strong_migrations: Catch unsafe migrations in development. https://github.com/ankane/strong_migrations.
- [10] Hao Li, Haoxiang Zhang, and Ahmed E. Hassan. 2025. The Rise of AI Teammates in Software Engineering (SE) 3.0: How Autonomous Coding Agents Are Reshaping Software Engineering. *arXiv:2507.15003*.
- [11] Lovable. 2026. Connect to Supabase. Lovable Documentation, <https://docs.lovable.dev/integrations/supabase>.
- [12] Nat McAleese, Rai Michael Pokorny, Juan Felipe Ceron Uribe, et al. 2024. LLM Critics Help Catch LLM Bugs. *arXiv:2407.00215*.
- [13] Neon. 2026. Branching — Neon Documentation. <https://neon.com/docs/introduction/branching>.
- [14] Shlomi Noach. 2016. gh-ost: GitHub’s online schema migration tool for MySQL. <https://github.com/github/gh-ost>.
- [15] Raja Parasuraman and Dietrich H. Manzey. 2010. Complacency and Bias in Human Use of Automation: An Attentional Integration. *Human Factors* 52, 3 (2010), 381–410.
- [16] Andrew Pavlo, Gustavo Angulo, Joy Arulraj, et al. 2017. Self-Driving Database Management Systems. In *Proc. CIDR*.
- [17] Replit. 2026. Development and production databases. Replit Documentation, <https://docs.replit.com/cloud-services/storage-and-databases/production-databases>.
- [18] Javier Ron, Benoit Baudry, and Martin Monperrus. 2026. N-Version Programming with Coding Agents. *arXiv:2606.20158* (2026).
- [19] Mrinank Sharma, Meg Tong, Tomasz Korbak, et al. 2024. Towards Understanding Sycophancy in Language Models. In *International Conference on Learning Representations (ICLR)*.
- [20] Simon Sharwood. 2025. Vibe coding service Replit deleted user’s production database, faked data, told fibs galore. The Register. https://www.theregister.com/2025/07/21/replit_saastr_vibe_coding_incident/
- [21] Tianneng Shi, Jingxuan He, Zhun Wang, et al. 2025. Progent: Securing AI Agents with Privilege Control. *arXiv:2504.11703* (2025).
- [22] Tae-Eun Song. 2026. Cross-Context Review: Improving LLM Output Quality by Separating Production and Review Sessions. *arXiv:2603.12123* (2026).
- [23] Ningzhi Tang, Chaoran Chen, Gelei Xu, et al. 2026. How Coding Agents Fail Their Users: A Large-Scale Analysis of Developer-Agent Misalignment in 20,574 Real-World Sessions. *arXiv:2605.29442* (2026).
- [24] Richard Tibbetts. 2026. *Schema Migration Lifecycles at Scale: A Million-File Corpus Study*. Technical Report. Bulrush Labs. In preparation.
- [25] Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. Language Models Don’t Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [26] Dana Van Aken, Andrew Pavlo, Geoffrey J. Gordon, and Bohan Zhang. 2017. Automatic Database Management System Tuning Through Large-scale Machine Learning. In *Proceedings of the 2017 ACM International Conference on Management of Data (SIGMOD)*. 1009–1024.

- [27] Yuhao Wu, Franziska Roesner, Tadayoshi Kohno, Ning Zhang, and Umar Iqbal. 2025. IsolateGPT: An Execution Isolation Architecture for LLM-Based Agentic Systems. In *Proc. NDSS*.
- [28] Haoran Yu, Lifei Liu, Xiaochong Jiang, et al. 2026. Habituation at the Gate: Rising Approval and Declining Scrutiny in Human Review of AI Agent Code. arXiv:2606.22721 [cs.SE]
- [29] Xuanhe Zhou, Guoliang Li, Zhaoyan Sun, et al. 2024. D-Bot: Database Diagnosis System using Large Language Models. *PVLDB* 17, 10 (2024), 2514–2527.